

The Illusion of Value

In the spring of 2023, Oracle announced a major cloud partnership with OpenAI. Nvidia's stock surged as demand for its chips became inseparable from every serious AI venture. OpenAI's valuation climbed past \$80 billion, not primarily because its products were profitable, but because the companies surrounding it kept signaling that it was essential. No one lied, no press release was fraudulent, nor did any executive make a provably false claim. And yet, something was happening that deserves moral scrutiny. The perceived value of these companies was not simply being discovered by the market but constructed through a web of partnerships, investments, and public signals, each one reinforcing the others. My paper argues that businesses can act wrongly even when no individual firm directly deceives anyone. More specifically, I argue that participating in markets that sustain inflated beliefs about value is morally wrong, regardless of whether any single participant intended to mislead. The Oracle, OpenAI, and Nvidia triangle offers a clear and instructive example of exactly this kind of market, and it will serve as the running case throughout this paper.

Before defending this position, it is worth being precise about what I am and am not claiming. By "participants," I mean firms and decision-makers who actively shape and benefit from a market through investments, partnerships, and public signals. I am not claiming that every company involved in artificial intelligence is acting wrongly, nor that innovation in this space is without genuine value. My narrower claim is when firms participate in markets they have reason to know are sustaining beliefs about value that go beyond what the underlying technology currently justifies, they are doing something morally problematic, even if no single firm is lying. By "reason to know" I do not mean certainty, nor do I mean that firms must conduct their own independent audits of their own market impact. I mean something more modest and more demanding at the same time, that a firm has reason to know its signals are sustaining inflated

beliefs when two conditions are met. First, the firm understands how markets respond to signals of the kind it is sending. Second, the firm has internal access to information suggesting that the beliefs those signals generate are not supported by current performance or demonstrated value. Both conditions are clearly met in the Oracle, OpenAI, and Nvidia case. These are sophisticated market actors with legal, financial, and strategic teams whose entire function is to understand how their public moves will be received. They are not naive about the signals they send and each of them has internal knowledge of the uncertainties their public positioning does not reflect. Oracle's cloud partnership with OpenAI signals to the market that OpenAI's infrastructure needs are real and urgent. That signal increases demand for Nvidia chips, which reinforces the broader narrative that an AI boom is underway, which in turn makes OpenAI's valuation seem more justified. Each company benefits from the signals the others send and then on top of that NVIDIA invests in OpenAI, which OpenAI uses to purchase cloud computing from Oracle. Oracle then uses that revenue to purchase more AI chips from NVIDIA, reinforcing the partnership's revenue cycles. None of them have said nor needs to say anything false but the question I am pursuing is whether that is enough to make their participation morally acceptable. I will contend that it is not.

The most straightforward objection to calling this kind of market wrong is that no one is being lied to. If Oracle announces a partnership and Nvidia reports record chip demand, these are true statements. But again, the moral problem I am identifying does not depend on anyone saying anything false. It depends on what the overall system of signals produces in the minds of those who rely on them. When Oracle invests in OpenAI, the investment is not just a business decision but it functions as a public signal that OpenAI's technology is credible, scalable, and worth staking real resources on. When Nvidia's valuation rises in response, that rise itself becomes

another signal, one that suggests the underlying demand is real and growing. Each signal makes the next one more believable and the result is that value is not being discovered through these interactions but being constructed by them. When people make decisions that affect their financial security, their careers, or how they allocate their resources, they are entitled to make those decisions on the basis of accurate information to the extent that accurate information is available. I don't believe this to be a controversial claim and in fact we already accept it in other contexts. We require drug companies to disclose side effects, we require sellers of securities to disclose material risks, we hold advertisers to standards of truthfulness, and so on. The moral principle underlying all of these practices is the same, that people who stand to be significantly affected by a decision have a legitimate claim to the information that would allow them to make that decision well. When a market systematically produces signals that outrun the reality they claim to represent, it violates that claim, not by lying, but by creating conditions in which informed decision-making becomes structurally impossible for those on the outside. This matters morally because the people making decisions based on these signals, smaller investors, employees choosing where to work, companies deciding whether to build on these platforms, are not evaluating the technology on its merits. They are evaluating a story that the participants are collectively telling, and that the participants have reason to know is running ahead of what the technology currently justifies. Participating in that process and benefiting from it is wrong for the same reason that any practice is wrong when it predictably leads others to make important decisions based on a distorted picture of reality, even if no single participant drew the distortion directly themselves.

What makes this kind of participation not just morally problematic but specifically unfair is the unequal distribution of knowledge among those affected by it. The firms at the center of

this market, Oracle, OpenAI, and Nvidia, have access to information that outside participants do not. They know how uncertain the underlying technology remains, how much of their valuation depends on narrative rather than demonstrated profitability, and that the partnerships generating market confidence are also generating the very signals that justify further investment. A retail investor reading about Oracle's expanded cloud infrastructure for OpenAI does not have access to any of this. What that investor sees is a signal, a credible company staking real money on another credible company, and draws the reasonable conclusion that the technology is as valuable as the market suggests. The problem is that the signal does not communicate the full picture, but instead functions primarily to communicate confidence. When the people sending signals know more than the people receiving them, and when those signals are shaping decisions involving real money, real careers, and real resource allocation, the resulting market is not just misleading. It is unfair in a way that matters morally, because it denies outside participants the information they would need to make genuinely informed decisions. Fairness in markets does not require that everyone knows everything. It does require that the signals people rely on are not systematically shaped by those who benefit from others believing them. This does not mean participation in competitive markets is inherently wrong, or that firms must disclose every internal uncertainty to the public. It means that firms with reason to know their signals are running ahead of their actual performance have a specific obligation to say so, or at minimum to refrain from amplifying those signals further. The moral problem is not partnership itself but the combination of inside knowledge and outside silence.

The most serious challenge to this paper's argument comes from a well-established tradition in moral philosophy that holds wrongdoing must be traceable to individual agents with genuine intent. Manuel Velasquez, in "Why Corporations Are Not Morally Responsible for

Anything They Do," argues that corporations cannot be morally responsible for their actions because moral responsibility requires the kind of intentionality and free will that only individual human beings possess. A corporation, in this view, is not the kind of thing that can intend to mislead anyone, because intending requires a mind, and corporations do not have minds in the morally relevant sense. When a corporation acts, what is really happening is that individual human beings are making decisions, and it is those individuals who bear whatever moral responsibility is attached to the outcome. Jan Narveson makes a closely related argument in his treatment of collective responsibility, contending that for genuine moral blame to apply, there must be a morally faulty decision traceable to a specific agent. Where no such decision exists, blame simply has nowhere to land. Together, Velasquez and Narveson represent what we call the individualist tradition in moral responsibility: the view that systemic outcomes, however harmful, do not generate moral wrongdoing unless we can point to specific people who chose to bring them about. This is a serious position and not easily dismissed. It captures something deeply intuitive about how we normally think about blame. We do not hold the wind responsible for knocking over a tree, and we are rightly suspicious of moral frameworks that assign blame so broadly that no one in particular is actually responsible for anything.

Applied to the Oracle, OpenAI, and Nvidia triangle, this objection has real force. Oracle's executives made a business decision to expand cloud infrastructure. Nvidia's leadership responded to genuine market demand for its products. OpenAI's team pursued partnerships that would allow it to scale. None of these decisions, taken individually, looks like an act of deception or manipulation. Each firm was doing what firms are supposed to do, pursuing opportunities, communicating their activities publicly, and responding to market conditions. If we follow Velasquez and Narveson, we would need to identify a specific individual or firm that

made a morally faulty choice with intent to mislead before we could assign blame. And that individual or firm does not obviously exist in this case. The executives at these companies could reasonably say they believed in what they were building, that their partnerships reflected genuine strategic value, and that any inflation of market expectations was a product of investor enthusiasm rather than deliberate engineering. On the individualist account, that defense is not just convenient, it is philosophically sound. If no one intended to mislead, and no single decision was morally faulty, then the system that resulted, however it looks from the outside, cannot properly be called wrong.

Velasquez and Narveson are right that moral responsibility requires more than just a harmful outcome. A world in which blame attaches to anyone who happens to be near a bad result, regardless of what they chose or knew, would be a world where moral responsibility loses its meaning entirely. That much I am willing to concede. But their framework carries a hidden assumption that deserves scrutiny. It was built to handle cases where the unit of wrongdoing is a discrete act by an identifiable agent. It works well when we are asking whether a specific executive authorized a fraudulent report, or whether a specific firm dumped toxic waste knowingly. It works less well, and I want to suggest it fails entirely, when the wrongdoing is not located in any single act but in a pattern of participation that no individual firm designed and no individual decision fully explains. The Oracle, OpenAI, and Nvidia triangle is precisely this kind of case. We cannot point to a boardroom meeting where executives agreed to inflate market expectations. We cannot identify the moment when participation crossed from legitimate business activity into something morally problematic. And yet the system that their collective participation produces, one in which signals reinforce each other, valuations climb beyond what demonstrated performance justifies, and outside participants make consequential decisions based

on a picture of value that insiders have reason to know is incomplete, is not morally neutral just because no single firm drew it. Velasquez contends that corporations lack the intentionality required for moral responsibility but the firms in this case do not need collective intent for their participation to be wrong. What they need, and what they have, is reason to know. Each of these firms understands how markets respond to signals. Each understands that their partnerships generate confidence that travels beyond what their technology currently justifies. Each benefits from that confidence continuing. The individualist framework asks, “Who intended to cause this harm?” That is the wrong question for this kind of case. The right question, and the one I think should guide our moral assessment here is, “Who had reason to know their participation was sustaining a distorted picture of value, and chose to continue anyway?” When the question is framed that way, the moral responsibility does not disappear into the system. It attaches, with appropriate weight, to each firm that keeps participating with open eyes. Narveson insists that without a morally faulty decision, blame has nowhere to land but I would submit that choosing to continue in a system you have reason to know is misleading others is itself a morally faulty decision, even if it is not the decision that created the system in the first place.

A natural objection to this paper's argument is that no one is forced to participate in the AI market. Investors choose to buy in, employees choose where to work, and smaller companies choose which platforms to build on. If people make decisions voluntarily, the argument goes, they bear responsibility for the outcomes of those decisions, and the firms whose signals they relied on cannot be blamed for choices that were ultimately someone else's to make. This objection has intuitive appeal, but it conflates two things that should be kept separate. Whether participation is voluntary and whether it is fully informed. A decision can be both voluntary and based on a distorted picture of reality at the same time. The retail investor who chooses to buy

into an AI fund after reading about Oracle's partnership with OpenAI has made a free choice, but again that choice was shaped by signals that the participating firms had reason to know were running ahead of what the technology justifies. Voluntariness is a necessary condition for fair markets, but it is not sufficient. What fairness also requires is that the signals people rely on when making their voluntary choices are not systematically shaped by those who benefit from others believing them. When that condition is not met, the fact that no one held a gun to anyone's head does not dissolve the moral problem, it simply relocates it.

Perhaps the most tempting objection to my position is that the AI market is not built on nothing. OpenAI's models do represent genuine technological achievement, Nvidia's chips are powering real applications that did not exist five years ago, and Oracle's infrastructure is serving real customers with real needs. If the underlying technology is genuinely valuable, the objection goes, then the market signals surrounding it are not distortions at all. They are simply enthusiasm for something that deserves enthusiasm, and calling that enthusiasm morally problematic seems to punish firms for succeeding. This objection is worth taking seriously because it is partly right. I am not claiming that the technology is worthless, that the partnerships are purely theatrical, or that nothing of value is being built. My claim is more narrow and more specific than that. Real innovation and morally problematic market behavior are not mutually exclusive. A firm can be genuinely building something valuable and simultaneously participating in a system that sustains beliefs about that value which go beyond what current performance justifies. The question is not whether OpenAI has produced real technology. It is whether the signals generated by the Oracle, OpenAI, and Nvidia triangle accurately represent the degree of that value, or whether they are running ahead of it in ways that outside participants cannot easily detect. Genuine innovation does not retroactively justify misleading signals about how much that innovation is currently

worth. If it did, any firm with a real product could claim immunity from scrutiny about how it presents that product to the market, and the moral standard I am defending here would have no practical force at all.

I began this paper with the straightforward observation that no one lied. Oracle, OpenAI, and Nvidia made true statements, pursued legitimate partnerships, and responded to real market demand. Yet the system their collective participation produced sustained beliefs about value that outside participants had no reliable way to evaluate on their own. This is enough to make their participation morally wrong, not because any firm intended to deceive, but because each had reason to know it was contributing to a pattern that others were relying on to make consequential decisions.

The individualist tradition, represented by Velasquez and Narveson, holds that without a specific wrongdoer making a specific faulty choice, moral blame has nowhere to land. Their framework is compelling in cases of discrete wrongdoing, but it cannot account for situations where harm is structural and produced by a pattern of individually reasonable decisions that together create something no single participant designed. In those cases, the more important moral question is not who intended to cause this, but who had reason to know their continued participation was sustaining it.

The AI market is not unique in this respect. Any market that depends on signals generated by those who benefit from others believing them, without mechanisms that allow outside participants to independently evaluate those signals, carries the same problem. The AI market is simply the clearest and most consequential current example. Permissible participation in such a market does not require firms to stop forming partnerships or communicating their activities publicly. It requires that the signals firms send bear a closer relationship to what current

performance actually justifies, and that firms with reason to know their signals are outrunning that performance take active steps to say so.

Until that standard is met, participation in markets like this one is not morally neutral. It is a choice, made with open eyes, to benefit from a picture of value that others cannot fully see.

References

Narveson, Jan. "Collective Responsibility." *The Journal of Ethics*, vol. 6, no. 2, 2002, pp. 179-198.

Velasquez, Manuel G. "Why Corporations Are Not Morally Responsible for Anything They Do." *Business and Professional Ethics Journal*, vol. 2, no. 3, 1983, pp. 1-18.